

We know setting up your own environment is hard...
SO WE DID IT FOR YOU hackyeah.hardcounter.alejz.pl



SENTINEL INTERLOCK



AI AGENTS ARE QUICK, BUT...

Mistakes are faster.
How do we make sure they are always compliant and aligned with our goals?

AI agents no longer just answer questions – they act. They access systems, call APIs and modify data autonomously. But most security controls still judge actions one at a time. An authorized action can still be wrong for the task. Sentinel controls actions before execution and verifies the final result afterwards.

SAFE ACTION $\neq$ SAFE WORKFLOW

01

Is this action authorized?

Deterministic policy checks identity, permissions, allowed tools, data access, budgets and known attack patterns before execution.

02

Does this action still serve the assigned task?

Trajectory supervision evaluates the whole sequence of actions to detect goal drift, skipped controls and suspicious detours.

03

Did the correct result actually happen?

An independent verifier checks the real external state against the Task Contract. The agent saying “done” is not proof of success.

OUR SOLUTION

INTERCEPT

ALLOW · REDACT · REQUIRE
APPROVAL · BLOCK

Every proposed action is
checked before execution.

PERSIST

IMMUTABLE TRACE ·
SEQUENCING · RECOVERY

Every decision and effect
becomes durable evidence.

SUPERVISE

GOAL DRIFT · RISK ·
POSTCONDITIONS

We evaluate the full trajectory
and independently check the
final external state.

The control system lives outside the agent. The agent cannot disable it.

HYBRID DEFENSE

Fast checks first. Semantic reasoning where it matters.



DETERMINISTIC

Cheap, predictable and explainable. Runs before execution.

- Identity & permissions
- Tool and model allowlists
- PII & secret detection
- Egress restrictions
- Token & tool-call budgets
- Known exploit signatures
- Model source validation



SEMANTIC + TRAJECTORY

Used selectively when context matters – especially before high-impact actions.

- Prompt injection
- Goal drift
- Suspicious detours
- Skipped controls
- Context-dependent risk
- Whole-session behaviour

Semantic controls may tighten restrictions. They can never override a hard deny.

WE GAVE IT A BANKING AGENT

APPLICATION → DOCUMENTS → IDENTITY → SANCTIONS → RISK → DECISION

PROMPT INJECTION
HIDDEN IN A
DOCUMENT

BLOCKED

AGENT SKIPS
SANCTIONS
SCREENING

BLOCKED

PII SENT TO AN
UNAUTHORIZED
DESTINATION

REDACTED

AGENT REPORTS
SUCCESS, BUT
PERSISTED STATE IS
WRONG

FAILED

One synthetic KYC environment. The same governed path for normal actions, attacks and failures.

OUR TEAM

- 01 Aleksander Jeżowski
- 02 Mateusz Błażejowski
- 03 Mikołaj Matałowski
- 04 Krzysztof Adamczyk
- 05 Krzysztof Radzewicz



TRY TO BREAK IT



GitHub
[https://github.com/HardCounter/
Hack-Yeah-2026](https://github.com/HardCounter/Hack-Yeah-2026)

Working Environment
@ hackyeah.hardcounter.alejez.pl